EurekaZero: Supplementary Telemetry

Candidate movement precedes admission; retention and independent learning remain separate questions.

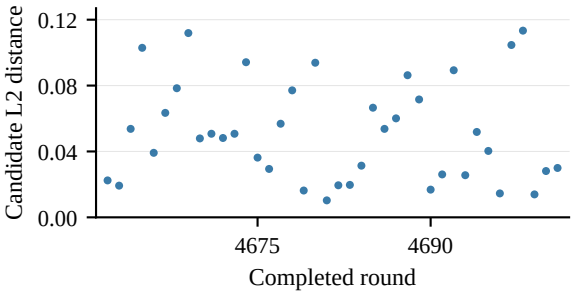

Figure S1: **Recorded candidate displacement varies across rounds.** All 40 complete retained summaries, rounds 4662–4701, captured at 00:55 UTC on 4 November 2026. Points show logged candidate parameter-space L2 displacement; no smoothing or filtering is applied. In the inspected loop, it includes warm-restart perturbation and optimization before admission. The journal records zero admissions in this window. This does not isolate optimizer effects or establish a retained learning gain. Scope: the adaptive adapter admission record in the main manuscript.

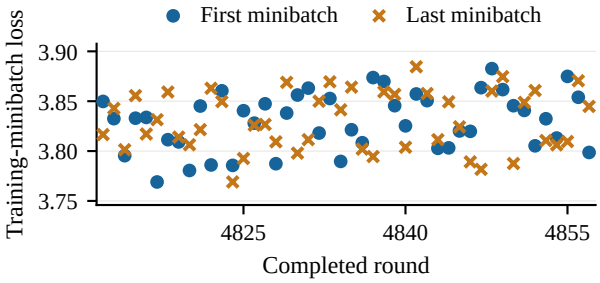

Figure S2: **First and last recorded minibatch losses use different batches.** All 46 complete retained summaries, rounds 4812–4857, captured at 05:53 UTC on 4 November 2026. The 92 points show logged first and last training-minibatch objectives, with two optimizer steps per candidate round. Batch-order hashes differ within every round; these are not matched before/after losses. No smoothing or filtering is applied. The window records zero admissions; these losses establish no retained learning gain. Scope: the adaptive adapter admission record in the main manuscript.